

Commissioned First-Use Evaluation of Metriplane 0.3.0 Across Six External Computing Environments

Protocol: MP-EXT-UX-001 v1.0

Evidence class: E2 commissioned external execution

Report version: 1.0.0

Report date: 2026-08-15

Compensation disclosure: All six primary evaluators were compensated. Payment was fixed and did not depend on success, favorable feedback, public praise, a GitHub star, or endorsement. PASS and FAIL were paid equally.

Claim boundary: This selected study does not establish organic adoption, customer use, company or institutional validation, market demand, production readiness, physical measurement accuracy, safety, generic robotics-data compatibility, generic ROS 2 compatibility, generic Isaac Sim compatibility, absence of other defects, or endorsement.

1. Executive summary

Six compensated external evaluators attempted the same frozen metriplane==0.3.0 first-use workflow on independently controlled environments. All six accepted primary records completed the bounded packaged workflow in their observed environments.

The accepted cohort contained four Native Linux environments, one macOS environment, one WSL2 Ubuntu environment, four x86_64 systems, two ARM64/AArch64 systems, four Python 3.12 environments, and two Python 3.13 environments. Primary evaluator compensation totaled USD 22 and EUR 57. Currencies are not converted, and platform fees are excluded.

Two additional commissioned records are retained as auxiliary evidence because their execution histories violated the frozen first-attempt rules. They are not included in the six-record denominator.

2. Evaluation question

Could a first-time external evaluator on an eligible pre-existing environment install Metriplane 0.3.0 from PyPI and complete the packaged camera-free incident, evidence-verification, and generated-regression workflow without maintainer troubleshooting before first-attempt evidence capture?

3. Evaluated artifact and method

- Package: metriplane==0.3.0 from PyPI, release tag v0.3.0.
- Resolved tag commit: e8ee6c63deae47bd450c5d6c7523d5bd699852a.
- Python boundary: 3.12 or 3.13. Core commands: metriplane --version, metriplane doctor, and metriplane demo.
- Eligible environments: Native Linux, macOS, or an already-provisioned WSL2 Ubuntu environment.

No camera, GPU, robot, Docker installation, or ROS installation was required. Browser opening was optional and outside the core result.

The frozen protocol required one measured attempt, no rehearsal, no rerun after a meaningful failure, stopping at the first meaningful failure, no maintainer or other troubleshooting before evidence preservation, fixed commands, and a pre-existing eligible environment. A valid product failure would remain a paid primary result. Protocol-invalid executions were retained separately as auxiliary evidence.

4. Recruitment and compensation

Evaluators were recruited through Fiverr and Freelancer for selected pre-existing, personally controlled environments. Compensation was outcome-neutral. Evaluator identity is not paired with compensation in this public package.

Slot	Platform	Compensation
T01	Fiverr	10 USD
T02	Fiverr	12 USD
T03	Freelancer	8 EUR
T04	Freelancer	20 EUR
T05	Freelancer	9 EUR
T06	Freelancer	20 EUR
Total	Currencies not converted	USD 22; EUR 57

5. Primary cohort

Slot	Environment	Architecture	Python	Core	Elapsed
T01	Ubuntu 24.04.4 LTS, Native Linux	x86_64	3.12.3	PASS	Not recorded
T02	Linux Mint 22.3, Native Linux	x86_64	3.12.3	PASS	Not recorded
T03	macOS 26.5.1	ARM64/AArch64	3.12.13	PASS	25 s exact
T04	Debian GNU/Linux 13, Native Linux	ARM64/AArch64	3.13.15	PASS	About 30 s
T05	Ubuntu 24.04.3 under WSL2	x86_64	3.12.3	PASS	About 7 min 18 s
T06	Fedora Linux 42, Native Linux	x86_64	3.13.11	PASS	About 2 min

6. Bounded results

- Core completion: 6 of 6 accepted primary records completed the bounded packaged workflow in their observed environments.
- Environment count: 4 Native Linux, 1 macOS, and 1 WSL2 Ubuntu.
- Architecture count: 4 x86_64 and 2 ARM64/AArch64. Python count: 4 Python 3.12 and 2 Python 3.13.
- Compensation totals: USD 22 and EUR 57, kept separate. Elapsed-time coverage: 1 exact, 3 approximate, and 2 not recorded.
- Assistance coverage: Five of six primary records explicitly captured assistance before evidence preservation, and all five recorded NO. T02 remains NOT_RECORDED.
- First-confusion coverage: Five of six primary records captured the field. Four recorded no confusion, one recorded WSL2 report-access confusion, and T02 remains NOT_RECORDED.

T02 limitation: T02 establishes successful installation, readiness, and bundled demo completion. It does not separately establish every later structured sub-result. The affected values remain NOT_RECORDED in public_results.csv.

AI-assistance use was not separately captured as a structured public field. No value was inferred later.

7. Missing data and methodology notes

No missing value was inferred, and no evaluator was rerun to fill a gap. The complete machine-readable missing-field lists are in `public_results.csv`.

Slot	Methodology note
T01	Environment disclosure. Pre-screening stated Ubuntu 22.04.4. The preserved one-time execution used an eligible pre-existing Ubuntu 24.04.4 environment. No rerun occurred.
T02	Missing structured fields. The public record predates the later structured template. Several later sub-results, assistance fields, and timing were not reconstructed.
T03	Prior repository exposure. The evaluator had viewed the repository but had not previously installed or executed Metriplane. A timing wrapper would not have stopped on failure, but no command failed.
T04	Public transcription detail. The copied public installation command appears to omit a closing quotation mark. Successful pinned package resolution establishes the intended command.
T05	WSL2 usability note. Core execution remained valid. Report-access and stable-filename observations led to a documentation-only response.
T06	First-use confirmation timing. The strongest first-ever-install and run confirmation was captured after the measured run.

Data treatment

The public package keeps each unknown field as `NOT_RECORDED` rather than treating it as `NO`, `PASS`, or an empty value. The report headline counts use only fields directly established in the canonical primary table.

The visible report uses human-readable labels. Machine-readable records retain controlled values for `PASS`, `FAIL`, `YES`, `NO`, `NOT_RECORDED`, `NOT_APPLICABLE`, `EXACT`, and `APPROXIMATE`.

Public technical result links remain in the canonical CSV and source manifest. Evaluator handles, private crosswalks, internal project-management URLs, order identifiers, receipts, messages, and private source records are not part of this package.

8. Auxiliary commissioned records

Record	First-attempt status	Later result	Reason excluded
AUX-A	Invalid precondition	Follow-up completed	The required WSL2 and Python environment was provisioned during execution, and work continued after failed preconditions.
AUX-B	Failed	Second run completed	Execution continued after failure, written working-directory guidance was supplied, and a rerun was performed.

The later technical completions remain useful auxiliary evidence, but they are not protocol-compliant primary first-attempt passes.

9. Engineering response

The WSL2 observations led to GitHub issue #60, pull request #61, and merge commit 5e5396f0756b47ff02f6f0831ec62a44ea21c118.

The remediation changed documentation only. It added writable-directory guidance, explicit output-path guidance, Windows report-opening guidance, and stable-filename guidance. Runtime code, CLI behavior, the cell_truth_report.html filename, evidence verification, generated regression behavior, the v0.3.0 tag, and the original cohort remained unchanged. No post-fix run was added to the original denominator.

10. Limitations

The sample was small, selected, compensated, and not random. Results apply only to the observed environments and bounded packaged workflow.

- No physical robot, factory, camera, or physical measurement was evaluated.
- No third-party robotics source, ROS 2 integration, or Isaac Sim integration was evaluated.
- No safety case, certification, production deployment, or reliability claim was evaluated.
- Some fields were not recorded for every record. WSL2 browser dispatch remains environment-dependent. Other defects may remain.

Interpretation: A completed bounded workflow in six selected environments is not a population estimate, compatibility guarantee, adoption measure, or endorsement.

11. Claim boundaries

Permitted claim. Six compensated external evaluators produced six accepted primary records for the same frozen Metriplane 0.3.0 packaged first-use workflow. All six accepted records completed that bounded workflow in their observed environments.

Required disclosure. State that the cohort was selected, commissioned, and compensated. State that the result applies only to the bounded workflow and observed environments. Keep the two auxiliary records outside the six-record denominator.

Prohibited claims. Do not describe this result as organic adoption, customer use, company or institutional validation, market validation, production readiness, safety validation, robot or factory validation, broad platform support, generic ROS 2 compatibility, generic Isaac Sim compatibility, universal installation success, absence of defects, reliability certification, or endorsement.

12. Evidence and reproducibility

- public_results.csv is the canonical six-row primary table.
- auxiliary_records.csv is the canonical two-row auxiliary table.
- summary.json is generated from those two CSV files.
- SOURCE_MANIFEST.json maps material claims to public GitHub sources and local public data files.
- SHA256SUMS records the hashes of the committed public artifacts.

The public package contains cohort slots and public technical sources only. Private crosswalks, private source records, contracts, billing records, order records, messages, and internal project-management links are not included.

13. Conclusion

The six accepted commissioned records support only the bounded statement that the packaged Metriplane 0.3.0 workflow completed in the six observed environments under the recorded conditions. The WSL2 usability finding produced a documentation-only response. The result must not be upgraded into adoption, organizational validation, production readiness, safety, or general-platform claims.

Archive and DOI status: HOLD pending owner approval.

Public report package: docs/evaluations/commissioned-first-use-v030/